WaveCatBoost: Hybrid Air Quality Forecasting

Applied Time Series Analysis — Course Project

Group 8

Kumar Utkarsh
240150016

P Sai Deva Charan
240150024

M Charan Sree Teja
240150018

D Lalit Santosh
240150044

Problem and Approach

Why this is hard

- Air-quality forecasting needs **near real-time** answers.
- That pushes us toward **lighter** models (trees, short lags).
- The series is **non-stationary**: spikes, regime shifts, and changing variance.
- Light models struggle on the **raw** mix of spikes, diurnal cycles, and slow drift.

Paper approach

1. Split the series into scale-wise components (details + smooth) with Haar MODWT.
2. Forecast each component with CatBoost.
3. Reconstruct by summing bands; optionally add conformal intervals.

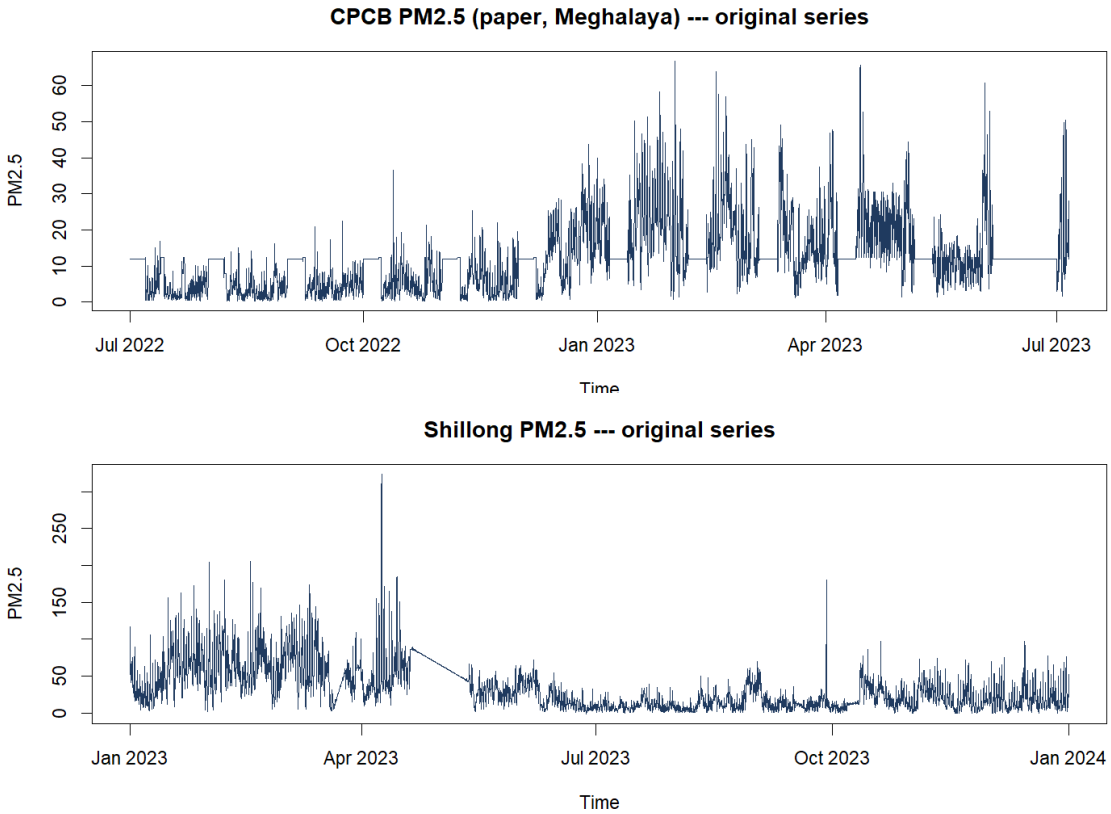
Datasets

Paper CSVs (Meghalaya, hourly)

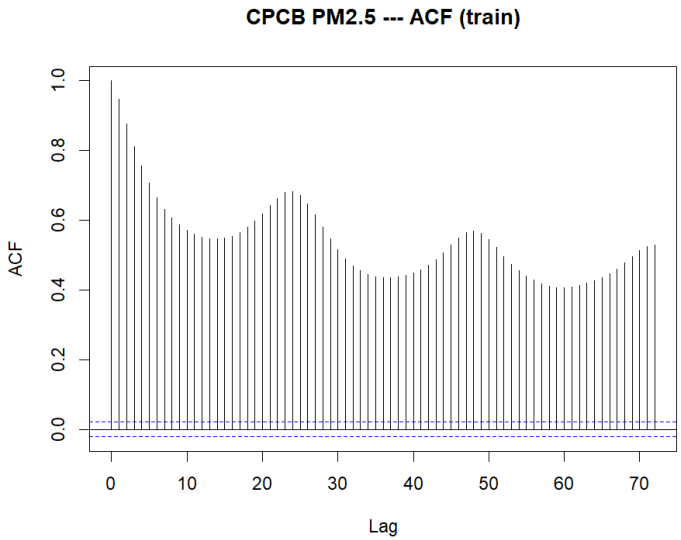
- **CPCB**, **ID1** from Borah et al.
- $n \approx 8880$ / 9505.
- NO_2, O_3, CO, SO_2, $PM_{2.5}$, PM_{10}.

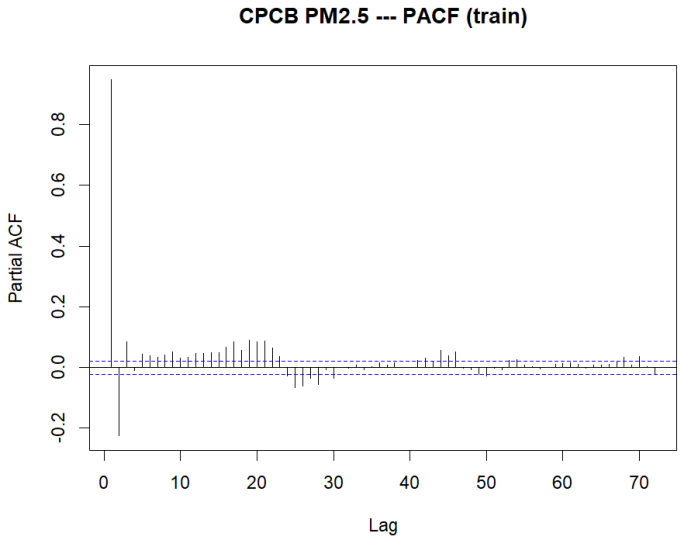
Extra CAAQMS (15-min $\rightarrow$ hourly)

- **Shillong:** Lumpyngngad.



Dataset diagnostics — ACF / PACF





ACF

Correlation of Y_t with Y_{t-k} for lags up to 72.

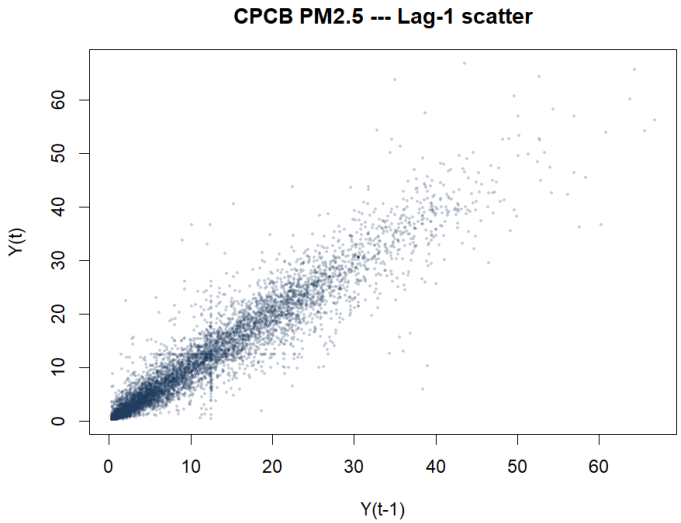
Stays high and waves every ~ 24 lags $\Rightarrow$ strong daily seasonality plus long memory (yesterday's level still matters).

PACF

"Extra" correlation after earlier lags are accounted for.

Huge spike at lag 1, then mostly small $\Rightarrow$ strongest short-term link is the previous hour (AR-style / lag-1 features). Lag 24 shows more in the ACF than as a big PACF spike.

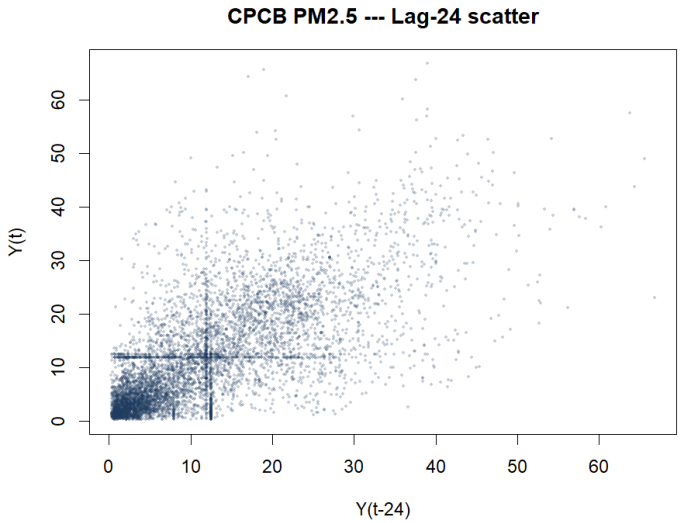
Dataset diagnostics — lag scatters



Lag-1 scatter

Y_t vs Y_{t-1}: tight diagonal.

Next hour is usually close to this hour $\Rightarrow$ high persistence / smooth hour-to-hour changes.

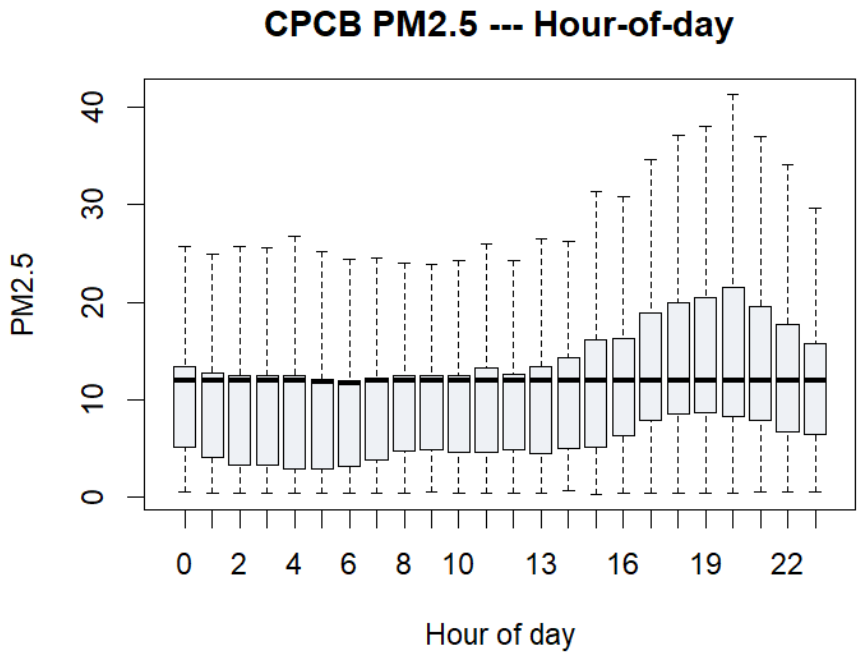


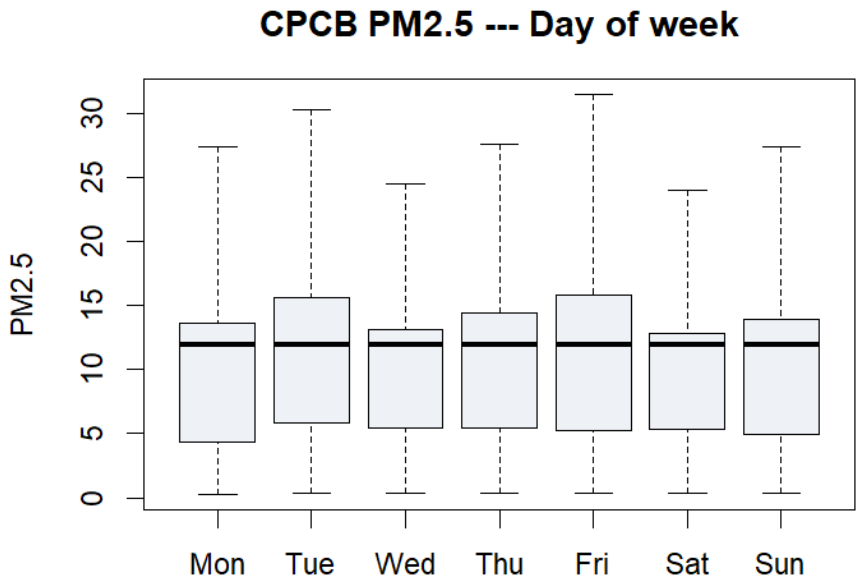
Lag-24 scatter

Y_t vs Y_{t-24}: still positive (same hour yesterday helps), but much looser than lag 1.

Daily cycle is real, but weaker than the immediate lag.

Seasonal structure — CPCB $PM_{2.5}$





- Clear diurnal pattern; milder weekday effect.
- S = **seasonality period** (hours in one cycle). Here $S{=}24$ = one day.
- Used for seasonal naive, SARIMA seasonal lag, and the MASE denominator.
- Also motivates lag features up to 24 h.

What we did before modelling

- **Hourly grid:** end-labelled hourly means (15-min CAAQMS $\rightarrow$ hour t).
- **Quality screen:** physical bounds, stuck sensors, spike flags.
- **Causal fill:** impute only with information available up to the cut time.
- **Chronological split:** train $\rightarrow$ val (14 d tuning) $\rightarrow$ test (last h hours); $h \in \{24, 168, 336, 744\}$. No random split.
- **Scale:** min–max to $[0, 1]$ fit on **train only**.
- Paper Meghalaya CSVs are already hourly $\Rightarrow$ skip re-aggregation / fill. Extra Shillong data get the full clean path into an isolated store.

Why decompose? Why not Fourier?

Why split the series

- Raw Y_t mixes noise, diurnal swings, and slow level shifts.
- One light model on the mix misses scale-specific structure.
- Decompose $\rightarrow$ forecast each band $\rightarrow$ add back $=$ richer dynamics without a heavy global model.

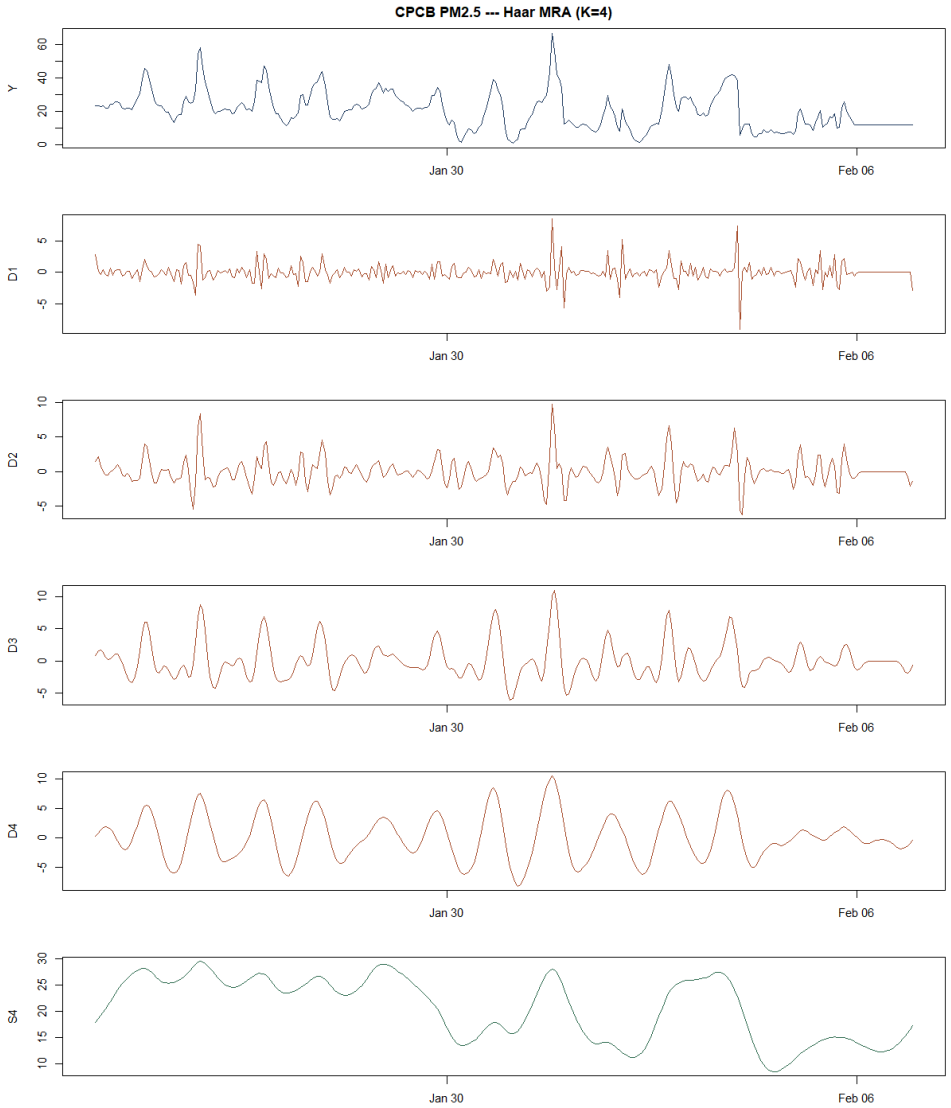
Why not Fourier

- Fourier uses **global** sinusoids over the whole record.
- Air quality is **non-stationary** — spikes and regime changes are local in time.
- Wavelets give a joint **time–scale** view; MODWT keeps every band length N (no downsampling).
- MODWT MRA is additive for Haar, D4, LA8, ...: $Y_t = \sum_k D_{k,t} + S_{K,t}$ (paper uses Haar).

MODWT: additive components

$$Y_t = \sum_{k=1}^{K} D_{k,t} + S_{K,t}$$

- $D_{k,t}$: detail at scale $\sim 2^k$ h.
- $S_{K,t}$: smooth after K levels.
- Paper: $K \approx \ln N$; we use K=4–5 ($\sim$ 16–32 h smooth).
- D_1: high-frequency component;
  D_k: lower-frequency component;
  S_K: smooth trend / level.



MODWT: mathematical intuition

Wavelet filter: $p_m \to D$ (details)
$\sum_m p_m = 0$ (measures change)

Scaling filter: $q_m \to S$ (smooth)
$\sum_m q_m = 1$ (measures average)

Filter condition (nonzero n):
$\sum_m p_m p_{m+2n} = \sum_m q_m q_{m+2n} = 0$

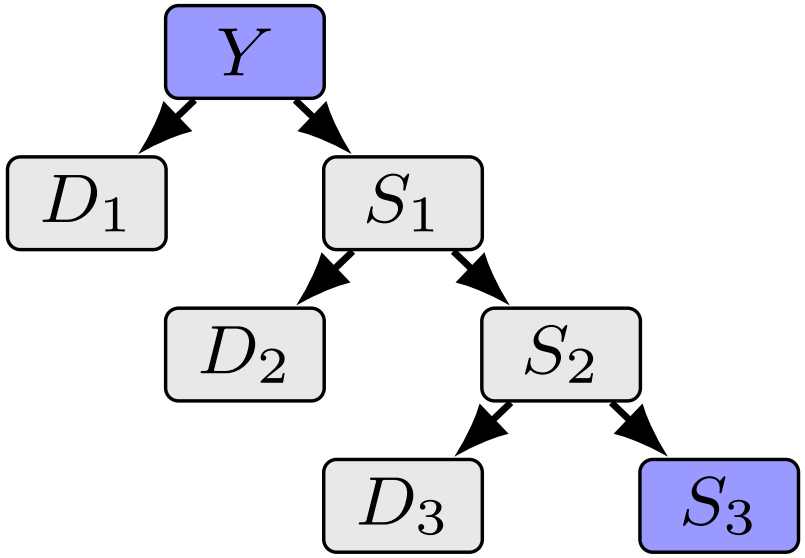
Recursion at level j ($S_{0,t} = Y_t$)

$D_{j,t} = \sum_m p_m \, S_{j-1,(t-2^{j-1}m) \bmod N}$

$S_{j,t} = \sum_m q_m \, S_{j-1,(t-2^{j-1}m) \bmod N}$

Haar: $q = (\frac{1}{2}, \frac{1}{2})$, $p = (\frac{1}{2}, -\frac{1}{2})$.

Tap spacing doubles each level.

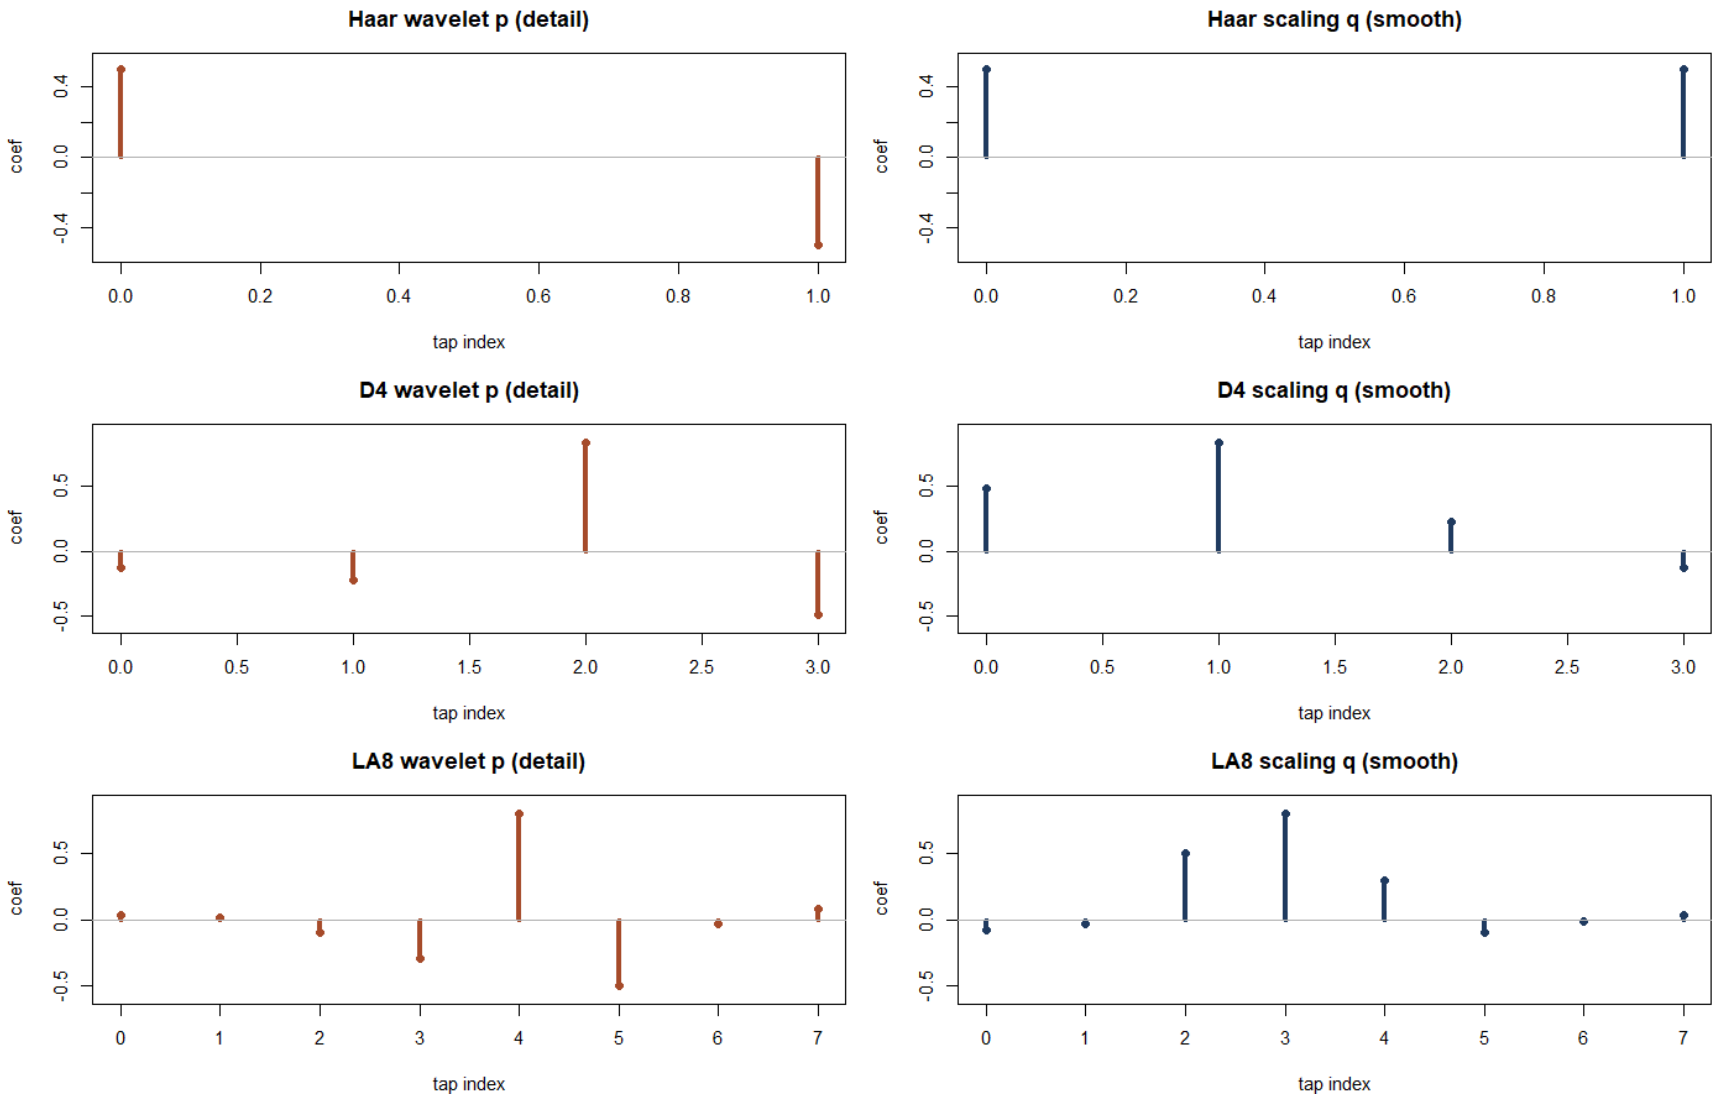


Only the smooth is split again.
$Y = D_1 + D_2 + D_3 + S_3$

Inverse (IMODWT): same filters, reversed

$S_{j-1,t} = \sum_m p_m D_{j,t+2^{j-1}m} + \sum_m q_m S_{j,t+2^{j-1}m}$

Filter coefficients



Which model to use?

Univariate ATS forecasting models

Naive baseline	Statistical models	Machine learning (ML)	Deep learning (DL)
Seasonal naive	SARIMA WaveSARIMA	**Raw lags** XGBoost LightGBM (LGBM) CatBoost **Wavelet hybrids** WaveXGB WaveLGB WaveCatBoost	**Examples** RNN TCN TFT Transformer

Error metrics

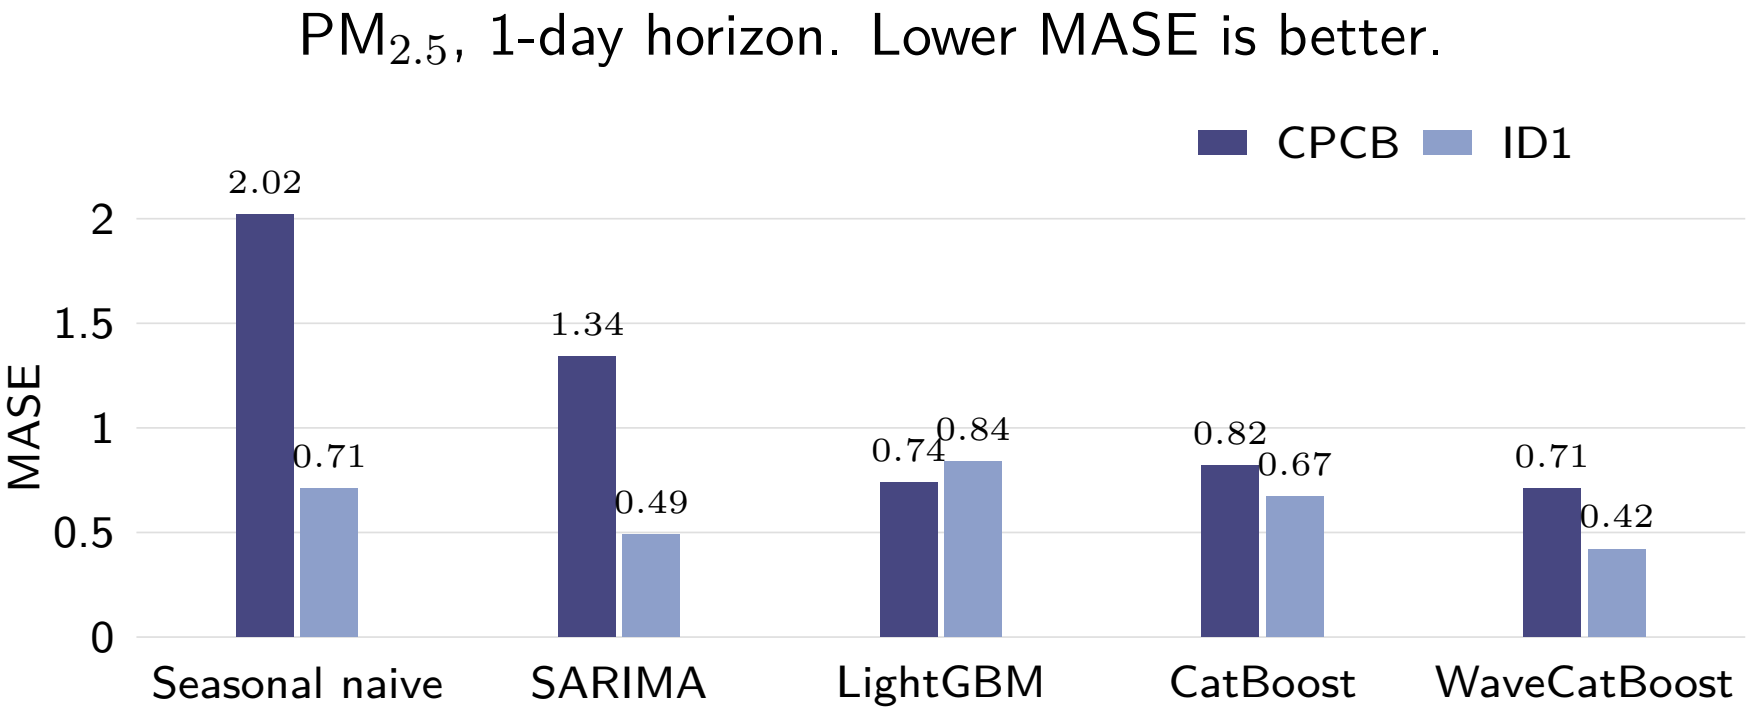
Mean absolute scaled error (MASE)

N is the training length, h is the forecast horizon, and $S = 24$ hours.

$$\mathrm{MASE} = \frac{N - S}{h} \frac{\sum_{t=N+1}^{N+h} |\hat{Y}_t - Y_t|}{\sum_{t=S+1}^{N} |Y_t - Y_{t-S}|}.$$

The denominator measures seasonal-naive error on the training data.

Model comparison



Sensor	Seasonal naive	SARIMA	LightGBM	CatBoost	WaveCatBoost
CPCB	2.02	1.34	0.74	0.82	0.71
ID1	0.71	0.49	0.84	0.67	0.42

CatBoost: From Trees to Ordered Boosting

1. Decision tree

$$f(x) = w_j \quad \text{if } x \in R_j$$

Each leaf predicts a value.

2. Gradient boosting

$$F_0(x) = c, \qquad r_i = y_i - F_{m-1}(x_i), \qquad h_m(x_i) \approx r_i$$

$$F_m(x) = F_{m-1}(x) + \eta h_m(x), \qquad \eta \text{ is the learning rate.}$$

$$h_m(x_i) \approx - \left. \frac{\partial L(y_i, F)}{\partial F} \right|_{F=F_{m-1}(x_i)}$$

3. Ordered boosting and symmetric trees

Ordered gradients use a model fitted on preceding observations. This reduces prediction shift.

Same split at each depth. 2^d leaves, so $d = 6 \Rightarrow 64$ leaves.

Our Hyperparameters vs Original Paper

The paper followed a predefined protocol. We used a limited grid search to select the lag window, tree depth and learning rate.

Parameter	Original paper	Our implementation
Lag window (p)	24	$24, 48 \rightarrow 24$
Wavelet depth (K)	$\lfloor \ln N \rfloor$, capped at 9	5
Tree depth	Default 6	$4, 6 \rightarrow 6$
Learning rate (η)	Automatic selection	$0.05, 0.10 \rightarrow 0.05$

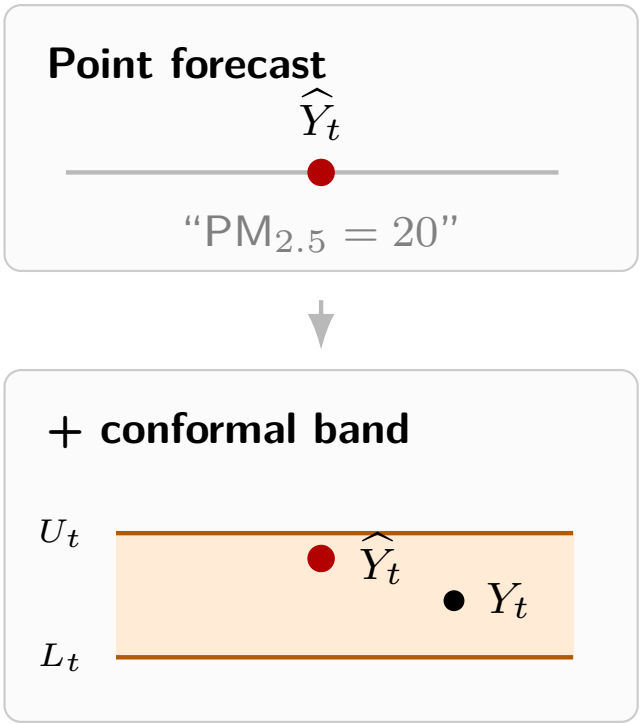
$$2\ (p) \times 2\ (\text{depth}) \times 2\ (\eta) = \boxed{8 \text{ combinations}}$$

$$\boxed{p = 24, \quad \text{depth} = 6, \quad \eta = 0.05}$$

Lag window: one daily cycle (24 h) versus two (48 h).
Tree depth: 4 versus 6. Learning rate: 0.05 versus 0.10.
Wavelet depth: $K = 5$, using the energy-based choice and practical limits. The CPCB theoretical depth is $K = 9$.

LightGBM selected the settings reused across the tree-based experiments. CatBoost was not independently tuned. The paper's core approach was retained.

Why conformal prediction?



- **Need:** alerts need a safe range, not one risky guess.
- **Recipe:** use recent $|Y - \widehat{Y}|$ to set width — large recent errors $\Rightarrow$ wider band.
- **No normal assumption:** AQ errors are spiky / quiet; conformal is distribution-free.
- **"95%":** about 95 of 100 hours, $Y_t \in [L_t, U_t]$ (coverage).

From $\widehat{Y}_t$ alone to a band $[L_t, U_t]$ with target coverage $1 - \alpha = 95\%$.

Conformal formula

Score $CS_t = \dfrac{|Y_t - \widehat{Y}_t|}{\widehat{U}_t}$

Quantile $CQ_t = Q_{1-\alpha}\big(CS_{t-\kappa+1:t}\big)$

Interval $PI_t = \widehat{Y}_t \pm CQ_t\,\widehat{U}_t$

Band edges: $L_t = \widehat{Y}_t - CQ_t\widehat{U}_t$, $U_t = \widehat{Y}_t + CQ_t\widehat{U}_t$ (clip $L_t \geq 0$).

- Y_t, $\widehat{Y}_t$: truth / WaveCatBoost.
- $\widehat{U}_t$: local error scale (next slide).
- CS_t: scaled |error| (nonconformity).
- α=0.05 $\Rightarrow$ 95% target.
- $\kappa \in \{24, 168, 336\}$ h (cal. window).

How to use it

1. Fit on train; residuals on a cal. window.
2. Build CS_t; pick κ for $\approx 95\%$ cal. coverage.
3. Test: band $\widehat{Y}_t \pm CQ_t\widehat{U}_t$.
4. Coverage = share of $Y_t \in [L_t, U_t]$.

How $\widehat{U}_t$ and CQ_t are computed

$\widehat{U}_t$ — local error scale

Residual $e_t = Y_t - \widehat{Y}_t$.

Default: **24-h rolling MAE**

$$\widehat{U}_t = \text{mean}\left(|e_{t-23}|, \ldots, |e_t|\right).$$

Large recent mistakes $\Rightarrow$ large $\widehat{U}_t$ $\Rightarrow$ wider band via $CQ_t\widehat{U}_t$.

Multi-step test (no new Y yet): freeze one robust $u^\star$ from calibration and reuse it.

$CQ_t = Q_{1-\alpha}(CS_{t-\kappa+1:t})$

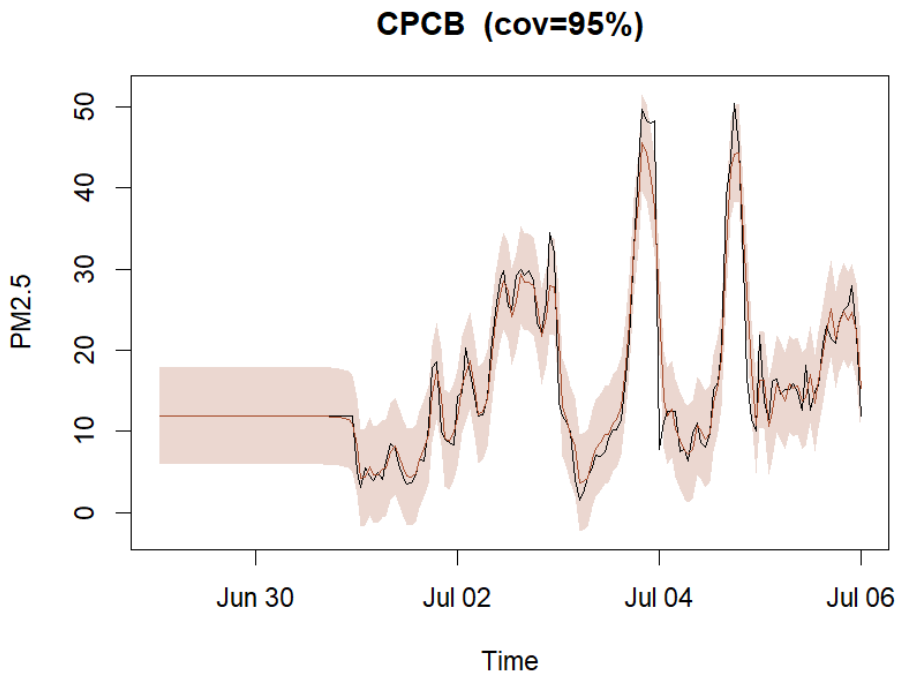
1. Take the last κ scores $CS_{t-\kappa+1}, \ldots, CS_t$.
2. Sort them.
3. Take the $(1-\alpha)$-quantile (α=0.05 $\Rightarrow$ $\approx$ 95th percentile).

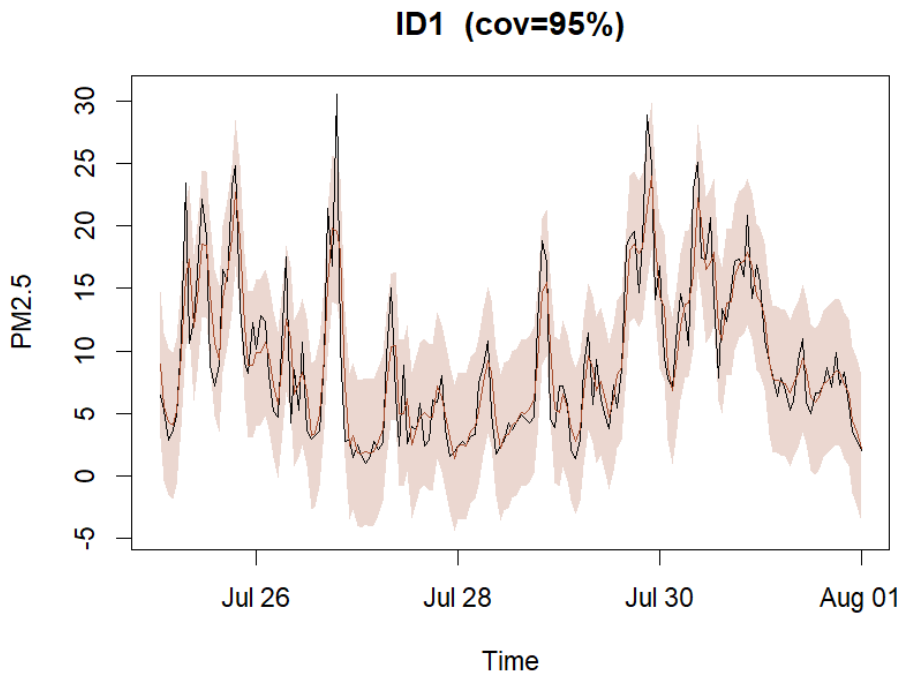
CQ_t is the cutoff for a "normal" scaled error.

Scores with $CS \leq CQ_t$ fall inside $\widehat{Y}_t \pm CQ_t\widehat{U}_t$.

Do the intervals work? (7-day test)

Purpose: letter-style 7-day $PM_{2.5}$ band — does truth sit inside the 95% window around WaveCatBoost?





- **Black** = actual; **orange** = WaveCatBoost; **shaded** = 95% residual PI $[L_t, U_t]$ (same colours as Rmd §8.2).
- Paper residual PI on this window: CPCB / ID1 coverage $\approx 95\%$ (widths ≈ 12.0 / 11.7).

7-day paper residual PI — summary

Sensor	Coverage	Mean width	Point MAE
CPCB	0.95	12.0	1.48
ID1	0.95	11.7	1.96
Shillong	0.95	19.3	3.56

Protocol

Rolling 1-step WaveCatBoost (letter notebook) + 95% quantile of $|Y - \widehat{Y}|$ on the plotted 7-day window ($\alpha = 0.05$).

Extensions — overview

- **Other wavelet filters:** Haar vs D4 vs LA8 (WaveCatBoost).
- **Fourier bands** vs Haar MODWT (WaveLGB).
- **Shillong A0–C3:** input / decomp designs on Lumpyngngad $PM_{2.5}$.

Haar vs D4 vs LA8 (WaveCatBoost, 1-day MASE)

Same WaveCatBoost head; only the MODWT filter (p_m, q_m) changes ($S_{0,t} = Y_t$, indices mod N):

$$D_{j,t} = \sum_{m=0}^{M-1} p_m \, S_{j-1,(t-2^{j-1}m) \bmod N}, \quad S_{j,t} = \sum_{m=0}^{M-1} q_m \, S_{j-1,(t-2^{j-1}m) \bmod N}.$$

Haar (M=2) $q = (\frac{1}{2}, \frac{1}{2})$, $p = (\frac{1}{2}, -\frac{1}{2})$
D4 (M=4) $q \approx (0.483,\ 0.837,\ 0.224,\ -0.129)$, $p \approx (-0.129,\ -0.224,\ 0.837,\ -0.483)$
LA8 (M=8) $q \approx (-0.076,\ -0.030,\ 0.498,\ 0.804,\ 0.298,\ -0.099,\ -0.013,\ 0.032)$
$p \approx (0.032,\ 0.013,\ -0.099,\ -0.298,\ 0.804,\ -0.498,\ -0.030,\ 0.076)$

Sensor	Pol.	Haar	D4	LA8	Sensor	Pol.	Haar	D4	LA8
CPCB	NO_2	0.541	**0.510**	0.529	ID1	NO_2	**1.252**	1.346	1.330
CPCB	$PM_{2.5}$	**0.727**	0.909	0.955	ID1	$PM_{2.5}$	**0.421**	0.457	0.502
CPCB	O_3	**1.014**	1.212	1.119	ID1	O_3	**0.935**	0.940	0.964
CPCB	CO	**0.885**	0.993	1.200	ID1	CO	**0.798**	1.001	1.170
CPCB	SO_2	6.232	6.041	**5.972**	ID1	SO_2	**1.320**	1.390	1.379
CPCB	PM_{10}	0.535	**0.524**	0.538	ID1	PM_{10}	**0.368**	0.445	0.462

Haar best on **9/12** cells; D4 on 2; LA8 on 1 — Haar stays the default.

Fourier vs Haar MODWT (WaveLGB, 1-day)

Octave FFT bands matched to Haar layout (J=5, p=24). Global sinusoids — no local $D_{1,t}$ (`extension_scores`, Rmd §10.1).

Sensor	Pollutant	Haar	Fourier
CPCB	NO_2	**0.518**	0.692
CPCB	$PM_{2.5}$	**1.031**	1.272
ID1	NO_2	**1.259**	1.438
ID1	$PM_{2.5}$	**0.360**	0.473

- Haar wins every scored cell.
- Fourier smears spikes across frequencies; lags still carry a little of the event.

Shillong: Experiment overview

Lumpyngngad $PM_{2.5}$: 24, 168, 336 and 744 hours

1. Input representation
Raw pollutant histories or causal Haar bands

↓

2. Information supplied
$PM_{2.5}$ / SO_2 and CO / Weather / Cross-band histories

↓

3. Model structure and learner
Separate or shared models; CatBoost, LightGBM, XGBoost

$PM_{2.5}$ forecasts: compare designs across forecast horizons

2023 split: Jan–Jul training / Aug–Sep validation / Oct–Dec test

Shillong A: Raw-input architectures

All designs also use hour-of-day and day-of-week inputs.

Design	Input features	Purpose
A0	Previous 24 hourly $PM_{2.5}$ values	$PM_{2.5}$ baseline
A1	A0 + humidity, wind speed, wind direction and rain	Weather information
A2	Previous 24 hourly values of $PM_{2.5}$, SO_2 and CO (72 pollutant inputs)	Other pollutants
A3	A2 + humidity, wind speed, wind direction and rain	Pollutants + weather

A0/A1 use $PM_{2.5}$ history. A2/A3 add SO_2 and CO histories.

Shillong B: Decomposed PM2.5 inputs

Causal Haar decomposition produces six components.

$$D_1,\ D_2,\ D_3,\ D_4,\ D_5,\ S_5$$

Design	Input and modelling setup
B0	Six $PM_{2.5}$ bands; one model per band
B1	Same bands; one shared model with a band identifier
B2	B0 + weather
B3	B1 + weather

B0/B2 fit six separate component models.
B1/B3 use one shared model and a band identifier.

Shillong C: Cross-component inputs

Can a band benefit from other bands' and pollutants' histories?

Design	Input features	Modelling setup
C1 (sep)	Histories of all six $PM_{2.5}$ bands (144 inputs)	Separate model per band
C1 (sh)	Same C1 information	One shared model
C2 (sep)	Histories of all six bands of $PM_{2.5}$, SO_2 and CO (432 inputs)	Separate model per band
C2 (sh)	Same C2 information	One shared model
C3 (sep)	C2 + weather	Separate model per band
C3 (sh)	C2 + weather	One shared model

Shillong: Main findings at Lumpyngngad

Multivariate pollutant inputs — strong benefit
SO_2 and CO histories improve $PM_{2.5}$ forecasts at every horizon.

Wavelet decomposition — no consistent gain
Raw pollutant histories generally perform better.

Weather inputs — horizon dependent
Little benefit at 24 h; often worse at longer horizons.

Cross-band information — mixed results
Helps at 24 and 744 h, but hurts at 168 h.

Architecture comparisons use LightGBM on the same test targets.

Shillong: Boosting-model comparison

Same four designs: A0, A3, B0 and C3 (sh). Four forecast horizons.
16 matched test comparisons. Lower mean rank is better.

Boosting model	Mean rank	Cases won (of 16)
CatBoost	1.625	9
LightGBM	2.000	4
XGBoost	2.375	3

A3 test MAE by horizon

Model	24 h	168 h	336 h	744 h
CatBoost	8.95	10.25	12.32	11.21
LightGBM	9.04	10.82	12.51	12.31
XGBoost	10.10	11.45	14.53	13.15

CatBoost wins the most matched comparisons among the three boosters.